

mod3

November 10, 2018

0.1 CBA Marathon Courses Fall 2018

1 Data Analysis and Visualization Using Python

1.1 Module 3 - Pandas

1.1.1 Dr. Mohammad AlMarzouq

2 Introduction to Pandas

- It is a python [3rd party library](#)
- Used for data analysis and visualization
- Part of Anaconda python distribution
- Best used with Jupyter notebook, can be used with regular python programs
- Main feature is the Data Frame

```
In [1]: # Load the pandas library to let python know you will use it
import pandas as pd
```

3 What is a Data Frame?

- Its a data structure, like lists and dictionaries
- Consists of rows and columns, similar to SQL tables and excel spreadsheets
- Columns are attributes or variables
- Rows are records or single observations
- Operations are typically performed on columns
- Has both numeric and named indexing

4 Tidy Data

- Standard form of organizing data in dataframe such that:
 - Each variable forms a column
 - Each row forms a row
 - Each table is an observational unit (level of analysis)
- Required reading: [Tidy Data, by Hadley Wickham](#)

5 Untidy/messy data can have

- Missing column names
- Aggregate/duplicate data
- Values for different variables in same column
- Data that should be on multiple tables is fitted into a single table
- ...more

6 Example of untidy/messy data

ID	Variable Name	Value
1	width	10
1	length	12
1	color	red
2	width	3
2	length	5
2	color	green

7 Example of tidy data

ID	width	length	color
1	10	12	red
2	3	5	green

8 Importance of Tidy Data

- Most tools we will use assume that data is tidy
- Collected data is likely messy or non-tidy, need to learn how to reshape it
- We will start with tidy data in our analysis
 - Will learn how to reshape when we start data collection

9 Loading data into a data frame

- Data is usually loaded/collected from an external source, like an api, website, csv, or excel file, and might not be well organized.
 - We will start with well organized and behaved data to get to know how to use Pandas
- Download the weather data set from [vega-dataset](#) (**right click and save as**)
- Place it in the same directory as the jupyter notebook you are working on

```
In [5]: # load the data using pandas library
        # do you remember what was pd?
```

```
pd.read_csv("weather.csv")
```

```
# Jupyter notebook tip:  
# type: pd.  
# then hit tab, see what happens  
# try also: pd.read_ (then hit tab)
```

```
Out[5]:
```

	location	date	precipitation	temp_max	temp_min	wind	\
0	Seattle	2012-01-01 00:00	0.0	12.8	5.0	4.7	
1	Seattle	2012-01-02 00:00	10.9	10.6	2.8	4.5	
2	Seattle	2012-01-03 00:00	0.8	11.7	7.2	2.3	
3	Seattle	2012-01-04 00:00	20.3	12.2	5.6	4.7	
4	Seattle	2012-01-05 00:00	1.3	8.9	2.8	6.1	
5	Seattle	2012-01-06 00:00	2.5	4.4	2.2	2.2	
6	Seattle	2012-01-07 00:00	0.0	7.2	2.8	2.3	
7	Seattle	2012-01-08 00:00	0.0	10.0	2.8	2.0	
8	Seattle	2012-01-09 00:00	4.3	9.4	5.0	3.4	
9	Seattle	2012-01-10 00:00	1.0	6.1	0.6	3.4	
10	Seattle	2012-01-11 00:00	0.0	6.1	-1.1	5.1	
11	Seattle	2012-01-12 00:00	0.0	6.1	-1.7	1.9	
12	Seattle	2012-01-13 00:00	0.0	5.0	-2.8	1.3	
13	Seattle	2012-01-14 00:00	4.1	4.4	0.6	5.3	
14	Seattle	2012-01-15 00:00	5.3	1.1	-3.3	3.2	
15	Seattle	2012-01-16 00:00	2.5	1.7	-2.8	5.0	
16	Seattle	2012-01-17 00:00	8.1	3.3	0.0	5.6	
17	Seattle	2012-01-18 00:00	19.8	0.0	-2.8	5.0	
18	Seattle	2012-01-19 00:00	15.2	-1.1	-2.8	1.6	
19	Seattle	2012-01-20 00:00	13.5	7.2	-1.1	2.3	
20	Seattle	2012-01-21 00:00	3.0	8.3	3.3	8.2	
21	Seattle	2012-01-22 00:00	6.1	6.7	2.2	4.8	
22	Seattle	2012-01-23 00:00	0.0	8.3	1.1	3.6	
23	Seattle	2012-01-24 00:00	8.6	10.0	2.2	5.1	
24	Seattle	2012-01-25 00:00	8.1	8.9	4.4	5.4	
25	Seattle	2012-01-26 00:00	4.8	8.9	1.1	4.8	
26	Seattle	2012-01-27 00:00	0.0	6.7	-2.2	1.4	
27	Seattle	2012-01-28 00:00	0.0	6.7	0.6	2.2	
28	Seattle	2012-01-29 00:00	27.7	9.4	3.9	4.5	
29	Seattle	2012-01-30 00:00	3.6	8.3	6.1	5.1	
...	
2892	New York	2015-12-02 00:00	3.0	13.9	8.3	2.0	
2893	New York	2015-12-03 00:00	0.0	13.3	7.2	7.2	
2894	New York	2015-12-04 00:00	0.0	11.7	5.0	4.7	
2895	New York	2015-12-05 00:00	0.0	11.7	1.7	2.4	
2896	New York	2015-12-06 00:00	0.0	10.6	3.3	2.9	
2897	New York	2015-12-07 00:00	0.0	12.8	4.4	3.4	
2898	New York	2015-12-08 00:00	0.0	10.6	4.4	3.5	
2899	New York	2015-12-09 00:00	0.0	12.8	1.1	3.4	
2900	New York	2015-12-10 00:00	0.0	15.0	8.9	3.0	

2901	New York	2015-12-11 00:00	0.0	14.4	7.8	2.7
2902	New York	2015-12-12 00:00	0.0	17.8	9.4	1.9
2903	New York	2015-12-13 00:00	0.0	21.1	11.7	3.1
2904	New York	2015-12-14 00:00	9.1	16.1	11.7	4.8
2905	New York	2015-12-15 00:00	2.3	17.8	11.7	8.2
2906	New York	2015-12-16 00:00	1.3	11.7	7.2	4.1
2907	New York	2015-12-17 00:00	29.7	15.0	10.0	4.1
2908	New York	2015-12-18 00:00	0.3	14.4	3.9	6.1
2909	New York	2015-12-19 00:00	0.0	5.0	2.2	9.0
2910	New York	2015-12-20 00:00	0.0	6.7	1.7	5.1
2911	New York	2015-12-21 00:00	0.0	12.8	3.3	5.3
2912	New York	2015-12-22 00:00	4.8	15.6	11.1	3.8
2913	New York	2015-12-23 00:00	29.5	17.2	8.9	4.5
2914	New York	2015-12-24 00:00	0.5	20.6	13.9	4.9
2915	New York	2015-12-25 00:00	2.5	17.8	11.1	0.9
2916	New York	2015-12-26 00:00	0.3	15.6	9.4	4.8
2917	New York	2015-12-27 00:00	2.0	17.2	8.9	5.5
2918	New York	2015-12-28 00:00	1.3	8.9	1.7	6.3
2919	New York	2015-12-29 00:00	16.8	9.4	1.1	5.3
2920	New York	2015-12-30 00:00	9.4	10.6	5.0	3.0
2921	New York	2015-12-31 00:00	1.5	11.1	6.1	5.5

	weather
0	drizzle
1	rain
2	rain
3	rain
4	rain
5	rain
6	rain
7	sun
8	rain
9	rain
10	sun
11	sun
12	sun
13	snow
14	snow
15	snow
16	snow
17	snow
18	snow
19	snow
20	rain
21	rain
22	rain
23	rain
24	rain

```

25      rain
26  drizzle
27      rain
28      rain
29      rain
...      ...
2892    fog
2893    sun
2894    sun
2895    sun
2896    sun
2897  drizzle
2898    sun
2899    sun
2900  drizzle
2901  drizzle
2902    fog
2903  drizzle
2904    fog
2905    fog
2906    fog
2907    fog
2908    sun
2909    sun
2910    sun
2911    sun
2912    fog
2913    fog
2914    fog
2915    fog
2916  drizzle
2917    fog
2918    snow
2919    fog
2920    fog
2921    fog

```

```
[2922 rows x 7 columns]
```

10 Loading File From URL

If the CSV file is downloadable from a url, you can put the URL in place of the file name:

```
In [5]: my_df = pd.read_csv("https://raw.githubusercontent.com/vega/vega-datasets/gh-pages/data/seattle_weather.csv")
my_df
```

```
Out[5]:
```

	location	date	precipitation	temp_max	temp_min	wind	weather
0	Seattle	2012-01-01	0.0	12.8	5.0	4.7	drizzle

1	Seattle	2012-01-02	10.9	10.6	2.8	4.5	rain
2	Seattle	2012-01-03	0.8	11.7	7.2	2.3	rain
3	Seattle	2012-01-04	20.3	12.2	5.6	4.7	rain
4	Seattle	2012-01-05	1.3	8.9	2.8	6.1	rain
5	Seattle	2012-01-06	2.5	4.4	2.2	2.2	rain
6	Seattle	2012-01-07	0.0	7.2	2.8	2.3	rain
7	Seattle	2012-01-08	0.0	10.0	2.8	2.0	sun
8	Seattle	2012-01-09	4.3	9.4	5.0	3.4	rain
9	Seattle	2012-01-10	1.0	6.1	0.6	3.4	rain
10	Seattle	2012-01-11	0.0	6.1	-1.1	5.1	sun
11	Seattle	2012-01-12	0.0	6.1	-1.7	1.9	sun
12	Seattle	2012-01-13	0.0	5.0	-2.8	1.3	sun
13	Seattle	2012-01-14	4.1	4.4	0.6	5.3	snow
14	Seattle	2012-01-15	5.3	1.1	-3.3	3.2	snow
15	Seattle	2012-01-16	2.5	1.7	-2.8	5.0	snow
16	Seattle	2012-01-17	8.1	3.3	0.0	5.6	snow
17	Seattle	2012-01-18	19.8	0.0	-2.8	5.0	snow
18	Seattle	2012-01-19	15.2	-1.1	-2.8	1.6	snow
19	Seattle	2012-01-20	13.5	7.2	-1.1	2.3	snow
20	Seattle	2012-01-21	3.0	8.3	3.3	8.2	rain
21	Seattle	2012-01-22	6.1	6.7	2.2	4.8	rain
22	Seattle	2012-01-23	0.0	8.3	1.1	3.6	rain
23	Seattle	2012-01-24	8.6	10.0	2.2	5.1	rain
24	Seattle	2012-01-25	8.1	8.9	4.4	5.4	rain
25	Seattle	2012-01-26	4.8	8.9	1.1	4.8	rain
26	Seattle	2012-01-27	0.0	6.7	-2.2	1.4	drizzle
27	Seattle	2012-01-28	0.0	6.7	0.6	2.2	rain
28	Seattle	2012-01-29	27.7	9.4	3.9	4.5	rain
29	Seattle	2012-01-30	3.6	8.3	6.1	5.1	rain
...
2892	New York	2015-12-02	3.0	13.9	8.3	2.0	fog
2893	New York	2015-12-03	0.0	13.3	7.2	7.2	sun
2894	New York	2015-12-04	0.0	11.7	5.0	4.7	sun
2895	New York	2015-12-05	0.0	11.7	1.7	2.4	sun
2896	New York	2015-12-06	0.0	10.6	3.3	2.9	sun
2897	New York	2015-12-07	0.0	12.8	4.4	3.4	drizzle
2898	New York	2015-12-08	0.0	10.6	4.4	3.5	sun
2899	New York	2015-12-09	0.0	12.8	1.1	3.4	sun
2900	New York	2015-12-10	0.0	15.0	8.9	3.0	drizzle
2901	New York	2015-12-11	0.0	14.4	7.8	2.7	drizzle
2902	New York	2015-12-12	0.0	17.8	9.4	1.9	fog
2903	New York	2015-12-13	0.0	21.1	11.7	3.1	drizzle
2904	New York	2015-12-14	9.1	16.1	11.7	4.8	fog
2905	New York	2015-12-15	2.3	17.8	11.7	8.2	fog
2906	New York	2015-12-16	1.3	11.7	7.2	4.1	fog
2907	New York	2015-12-17	29.7	15.0	10.0	4.1	fog
2908	New York	2015-12-18	0.3	14.4	3.9	6.1	sun
2909	New York	2015-12-19	0.0	5.0	2.2	9.0	sun

2910	New York	2015-12-20	0.0	6.7	1.7	5.1	sun
2911	New York	2015-12-21	0.0	12.8	3.3	5.3	sun
2912	New York	2015-12-22	4.8	15.6	11.1	3.8	fog
2913	New York	2015-12-23	29.5	17.2	8.9	4.5	fog
2914	New York	2015-12-24	0.5	20.6	13.9	4.9	fog
2915	New York	2015-12-25	2.5	17.8	11.1	0.9	fog
2916	New York	2015-12-26	0.3	15.6	9.4	4.8	drizzle
2917	New York	2015-12-27	2.0	17.2	8.9	5.5	fog
2918	New York	2015-12-28	1.3	8.9	1.7	6.3	snow
2919	New York	2015-12-29	16.8	9.4	1.1	5.3	fog
2920	New York	2015-12-30	9.4	10.6	5.0	3.0	fog
2921	New York	2015-12-31	1.5	11.1	6.1	5.5	fog

[2922 rows x 7 columns]

11 Now it is your turn

Download [airport.csv](#) then load it into the notebook

Remember: Right click on the link and select **save target as**

In []:

12 File Types

- CSV: Comma Separated Values [example](#)
 - Use `pd.read_csv`
- JSON: Javascript Object Notation [example](#)
 - Use `pd.read_json`
- Excel: Microsoft Excel File
 - Use `pd.read_excel`
- Others, type `pd.read_` then hit Tab to see a list

```
In [ ]: # Your turn
        # Load: https://github.com/vega/vega-datasets/raw/gh-pages/data/cars.json
        # Into: cars_df
```

13 How to work with the data?

- You must place it in a variable so you can refer to it
- The current data was displayed and not assigned to a variable, so you cannot use it
- Assign it to a variable named **my_df**

```
In [3]: my_df = pd.read_csv("weather.csv")
```

```
In [ ]: # Your turn: Load airports.csv into airports_df
```

14 Let us discover how the data looks like

We examine the top and bottom records of the dataframe to get an idea of what the data looks like

```
In [9]: my_df.head()
```

```
Out[9]:
```

	location	date	precipitation	temp_max	temp_min	wind	weather
0	Seattle	2012-01-01 00:00	0.0	12.8	5.0	4.7	drizzle
1	Seattle	2012-01-02 00:00	10.9	10.6	2.8	4.5	rain
2	Seattle	2012-01-03 00:00	0.8	11.7	7.2	2.3	rain
3	Seattle	2012-01-04 00:00	20.3	12.2	5.6	4.7	rain
4	Seattle	2012-01-05 00:00	1.3	8.9	2.8	6.1	rain

```
In [ ]: # You can pass a number in the head() method to show more data
        # show 10 items (try it)
```

```
        # do the same for airports_df
```

```
In [13]: # To know which columns are available use the columns attribute
         my_df.columns
```

```
Out[13]: Index(['location', 'date', 'precipitation', 'temp_max', 'temp_min', 'wind',
               'weather'],
              dtype='object')
```

```
In [ ]: # Your turn: explore the columns for airports_df
```

15 Data types

- Each **column** will have its own data type
- Remember, variables will be in columns
- Observations in rows
- Use dtypes attribute of to discover columns and datatypes
- **OOP**: What is the difference between a *function*, a *method*, an *attribute*, and a *variable*?

```
In [18]: my_df.dtypes
```

```
Out[18]: location      object
         date          object
         precipitation  float64
         temp_max      float64
         temp_min      float64
         wind          float64
         weather       object
         dtype: object
```

```
In [ ]: # Your turn: Find out the data types for the airports_df column
```



```
In [18]: # Pandas uses data types provided by numpy
# load numpy
import numpy as np
```

```
# convert the column to datetime
my_df.date.astype(np.datetime64)
```

```
Out[18]: 0      2012-01-01
1      2012-01-02
2      2012-01-03
3      2012-01-04
4      2012-01-05
5      2012-01-06
6      2012-01-07
7      2012-01-08
8      2012-01-09
9      2012-01-10
10     2012-01-11
11     2012-01-12
12     2012-01-13
13     2012-01-14
14     2012-01-15
15     2012-01-16
16     2012-01-17
17     2012-01-18
18     2012-01-19
19     2012-01-20
20     2012-01-21
21     2012-01-22
22     2012-01-23
23     2012-01-24
24     2012-01-25
25     2012-01-26
26     2012-01-27
27     2012-01-28
28     2012-01-29
29     2012-01-30
...
2892   2015-12-02
2893   2015-12-03
2894   2015-12-04
2895   2015-12-05
2896   2015-12-06
2897   2015-12-07
2898   2015-12-08
2899   2015-12-09
2900   2015-12-10
2901   2015-12-11
```

```
2902    2015-12-12
2903    2015-12-13
2904    2015-12-14
2905    2015-12-15
2906    2015-12-16
2907    2015-12-17
2908    2015-12-18
2909    2015-12-19
2910    2015-12-20
2911    2015-12-21
2912    2015-12-22
2913    2015-12-23
2914    2015-12-24
2915    2015-12-25
2916    2015-12-26
2917    2015-12-27
2918    2015-12-28
2919    2015-12-29
2920    2015-12-30
2921    2015-12-31
Name: date, dtype: datetime64[ns]
```

```
In [5]: # an alternative way to do it is using
pd.to_datetime(my_df.date).head() # do you remember head method?
```

```
Out[5]: 0    2012-01-01
1    2012-01-02
2    2012-01-03
3    2012-01-04
4    2012-01-05
Name: date, dtype: datetime64[ns]
```

```
In [6]: # now let us examine the date column
my_df.date.head()
```

```
# why is it still of type object?
# How to fix it?
```

```
Out[6]: 0    2012-01-01 00:00
1    2012-01-02 00:00
2    2012-01-03 00:00
3    2012-01-04 00:00
4    2012-01-05 00:00
Name: date, dtype: object
```

```
In [ ]: # just like the dataframe, the command creates a copy
# but does not store it
# We need to replace the old date column with the new one
my_df.date = my_df.date.astype(np.datetime64)
```

```
In [23]: # check the types
my_df.dtypes
```

```
Out[23]: location          object
date          datetime64[ns]
precipitation    float64
temp_max        float64
temp_min        float64
wind            float64
weather         object
dtype: object
```

```
In [2]: # Your turn: examine the airports_df dataframe
# are there any date columns that you can convert?
# Check then numeric columns, what should their data type be?
```

16 Data Types and Variable Types

Variable Type	Data Type
Continuous	float64, datetime64[ns]
Discrete	int64
Ordinal	int64, category
Nominal	int64, object, category
Categorical	int64, object, category

17 Why convert an object column into a date column?

- As you will find out later, pandas can do more fancy things if it knows the column is a date
- For example:
- Sort
- Filter based on date range
- Date arithmetic
- Always make sure date/time columns have the correct data type

18 Indexing Columns

- Using square brackets []
- Using dot notation .

```
In [7]: # a single column is known as a series
my_df['location']
```

```
Out[7]: 0      Seattle
1      Seattle
2      Seattle
```

3	Seattle
4	Seattle
5	Seattle
6	Seattle
7	Seattle
8	Seattle
9	Seattle
10	Seattle
11	Seattle
12	Seattle
13	Seattle
14	Seattle
15	Seattle
16	Seattle
17	Seattle
18	Seattle
19	Seattle
20	Seattle
21	Seattle
22	Seattle
23	Seattle
24	Seattle
25	Seattle
26	Seattle
27	Seattle
28	Seattle
29	Seattle
	...
2892	New York
2893	New York
2894	New York
2895	New York
2896	New York
2897	New York
2898	New York
2899	New York
2900	New York
2901	New York
2902	New York
2903	New York
2904	New York
2905	New York
2906	New York
2907	New York
2908	New York
2909	New York
2910	New York
2911	New York

```

2912    New York
2913    New York
2914    New York
2915    New York
2916    New York
2917    New York
2918    New York
2919    New York
2920    New York
2921    New York
Name: location, dtype: object

```

```

In [8]: # Some methods that work on Dataframes also work on Series
my_df['location'].head()

```

```

Out[8]: 0    Seattle
        1    Seattle
        2    Seattle
        3    Seattle
        4    Seattle
Name: location, dtype: object

```

```

In [13]: # Dot notation to access series
my_df.location.head()

```

```

Out[13]: 0    Seattle
        1    Seattle
        2    Seattle
        3    Seattle
        4    Seattle
Name: location, dtype: object

```

```

In [ ]: # Your turn: Try to index the columns for airports_df using square brackets and dot notation
# Use head() to get an idea of what the data is

```

```

In [14]: # Descriptive statistics
my_df['location'].describe()

```

```

Out[14]: count          2922
        unique           2
        top      New York
        freq          1461
        Name: location, dtype: object

```

```

In [15]: # works also on dataframe
my_df.describe()

```

```

Out[15]:      precipitation    temp_max    temp_min      wind
count    2922.000000    2922.000000    2922.000000    2922.000000

```

mean	2.944764	16.769131	8.612320	4.101129
std	7.695286	8.644596	7.511776	1.880791
min	0.000000	-7.700000	-16.000000	0.400000
25%	0.000000	10.000000	3.300000	2.700000
50%	0.000000	16.100000	8.900000	3.800000
75%	1.800000	23.900000	13.900000	5.100000
max	118.900000	37.800000	26.700000	16.200000

```
In [26]: # Different data types will have different descriptives
my_df['date'].describe()
```

```
Out[26]: count          2922
unique          1461
top      2013-06-05 00:00:00
freq              2
first    2012-01-01 00:00:00
last      2015-12-31 00:00:00
Name: date, dtype: object
```

```
In [21]: my_df.precipitation.describe()
```

```
Out[21]: count    2922.000000
mean         2.944764
std         7.695286
min         0.000000
25%         0.000000
50%         0.000000
75%         1.800000
max        118.900000
Name: precipitation, dtype: float64
```

```
In [ ]: # Your turn: Use describe() on airports_df
```

```
# Which columns are included in describe?
```

```
# Try it on the columns that were excluded:
```

```
# Why were these columns excluded?
```

19 You can also plot a dataframe


- Pandas will try to show it in the best way possible
- Plotting from dataframe is very simplistic and used for quick univariate exploration

```
In [ ]: my_df.plot()
```

```
In [29]: # You need to tell pandas that you want to display plots in the notebook
%matplotlib inline
```

```
In [31]: # now try it
my_df.plot()
```

```
Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x1147020b8>
```




20 Don't forget!

Always include in your notebook:

```
# this is the first cell in your notebook
import pandas as pd
%matplotlib inline # dont forget this
```

```
In [32]: # It's more meaningful to plot Series
my_df['precipitation'].plot()
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x1147250f0>
```



```
In [9]: # Remember plots show change from one observation to the next
my_df['wind'].plot()
```


```
# in some cases it might not be useful
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x10b33f710>
```

```
In [40]: # you can try a histogram
my_df['wind'].hist()
```

```
# Which is useful to know distributions
```

```
Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x116d89240>
```

21 How can you find out if percipitation is usually high in the year or low?

In []: *# Your turn:*

In [45]: *# Sometime pandas cannot plot it*
 my_df['location'].plot()

 TypeError

Traceback (most recent call last)

```
<ipython-input-45-e083b00ff51a> in <module>()
    1 # Sometime pandas cannot plot it
----> 2 my_df['location'].plot()

/Users/koutbo6/anaconda/lib/python3.5/site-packages/pandas/tools/plotting.py in __call__
3564         colormap=colormap, table=table, yerr=yerr,
3565         xerr=xerr, label=label, secondary_y=secondary_y,
-> 3566         **kwds)
3567     __call__.__doc__ = plot_series.__doc__
3568
```

```

/Users/koutbo6/anaconda/lib/python3.5/site-packages/pandas/tools/plotting.py in plot_s
2643         yerr=yerr, xerr=xerr,
2644         label=label, secondary_y=secondary_y,
-> 2645         **kwds)
2646
2647

```

```

/Users/koutbo6/anaconda/lib/python3.5/site-packages/pandas/tools/plotting.py in _plot(
2439     plot_obj = klass(data, subplots=subplots, ax=ax, kind=kind, **kwds)
2440
-> 2441     plot_obj.generate()
2442     plot_obj.draw()
2443     return plot_obj.result

```

```

/Users/koutbo6/anaconda/lib/python3.5/site-packages/pandas/tools/plotting.py in genera
1024     def generate(self):
1025         self._args_adjust()
-> 1026         self._compute_plot_data()
1027         self._setup_subplots()
1028         self._make_plot()

```

```

/Users/koutbo6/anaconda/lib/python3.5/site-packages/pandas/tools/plotting.py in _comput
1133     if is_empty:
1134         raise TypeError('Empty {0!r}: no numeric data to '
-> 1135                        'plot'.format(numeric_data.__class__.__name__))
1136
1137     self.data = numeric_data

```

TypeError: Empty 'DataFrame': no numeric data to plot

In []: *# Your turn: try to plot the columns in airport_df using either plot() or hist()*

What can you find out about the data?

In [49]: *# Such variables are usually categorical and you can get frequencies like so*
my_df['location'].value_counts()

```

Out[49]: New York      1461
Seattle      1461
Name: location, dtype: int64

```

```
In [3]: # Your turn: Examine the columns for airports_df
        # what would be the best columns to check frequencies for?
        # try it:

        # The best columns are:

        # The reason frequencies is best calculated on them is because:

        # What did you find out about your data?

        # try to plot value_count(), how do you do it? what do you get?
```

22 Are data frames immutable?

- Yes, however, all operations that change values will produce a copy and not change the original
- You have to use assignment to change columns or dataframes
- **So be careful!**

23 Missing Value (Nulls or Nans)

- Missing values are usually represented by:
 - The Python Null if the value doesn't exist
 - Numpy nan if the value is Not a Number (like zero division)
- Use `.dropna()` to remove rows with null values or `.fillna()` to replace the values

```
In [9]: # what are we doing here?
        cars_df = pd.read_json("https://github.com/vega/vega-datasets/raw/gh-pages/data/cars.json")

        # What are we doing here?
        len(cars_df)
```

Out[9]: 406

```
In [10]: # Why do we have a smaller value now?
         len(cars_df.dropna())
```

Out[10]: 392

```
In [11]: # let's look at the available columns
        cars_df.columns
```

```
Out[11]: Index(['Acceleration', 'Cylinders', 'Displacement', 'Horsepower',
               'Miles_per_Gallon', 'Name', 'Origin', 'Weight_in_lbs', 'Year'],
              dtype='object')
```

```
In [27]: # You can limit checking null values to specific columns
len(cars_df.dropna(subset=['Miles_per_Gallon', 'Horsepower', ]))
```

```
Out[27]: 392
```

```
In [28]: # condition can be that all are missing
len(cars_df.dropna(subset=['Miles_per_Gallon', 'Horsepower', ], how="all"))
```

```
Out[28]: 406
```

```
In [29]: # or at least one
len(cars_df.dropna(subset=['Miles_per_Gallon', 'Horsepower', ], how="any"))
```

```
Out[29]: 392
```

```
In [33]: # You can also choose to replace the nulls with a value
# like the mean or 0 depending on what makes more sense
cars_df.fillna(0).head()
```

```
Out[33]:
```

	Acceleration	Cylinders	Displacement	Horsepower	Miles_per_Gallon	\
0	12.0	8	307.0	130.0	18.0	
1	11.5	8	350.0	165.0	15.0	
2	11.0	8	318.0	150.0	18.0	
3	12.0	8	304.0	150.0	16.0	
4	10.5	8	302.0	140.0	17.0	

	Name	Origin	Weight_in_lbs	Year
0	chevrolet chevelle malibu	USA	3504	1970-01-01
1	buick skylark 320	USA	3693	1970-01-01
2	plymouth satellite	USA	3436	1970-01-01
3	amc rebel sst	USA	3433	1970-01-01
4	ford torino	USA	3449	1970-01-01

```
In [35]: # you can also specify which column gets which value using a dictionary
cars_df.fillna(
    {
        'Miles_per_Gallon':0, # put 0 instead of null in miled_per_gallon
        'Horsepower':cars_df.Horsepower.mean(), # put mean instead of null in horsepower
    }
).head()
```

```
Out[35]:
```

	Acceleration	Cylinders	Displacement	Horsepower	Miles_per_Gallon	\
0	12.0	8	307.0	130.0	18.0	
1	11.5	8	350.0	165.0	15.0	
2	11.0	8	318.0	150.0	18.0	
3	12.0	8	304.0	150.0	16.0	
4	10.5	8	302.0	140.0	17.0	

	Name	Origin	Weight_in_lbs	Year
0	chevrolet chevelle malibu	USA	3504	1970-01-01
1	buick skylark 320	USA	3693	1970-01-01
2	plymouth satellite	USA	3436	1970-01-01
3	amc rebel sst	USA	3433	1970-01-01
4	ford torino	USA	3449	1970-01-01

0	chevrolet chevelle malibu	USA	3504	1970-01-01
1	buick skylark 320	USA	3693	1970-01-01
2	plymouth satellite	USA	3436	1970-01-01
3	amc rebel sst	USA	3433	1970-01-01
4	ford torino	USA	3449	1970-01-01

24 Duplicates

- Search for any repeated values
- use duplicated()
 - Note that this looks at all the columns in the record
 - You can pass a list of column names to check duplication based on

In [28]: `my_df.duplicated()`

```
Out[28]: 0      False
         1      False
         2      False
         3      False
         4      False
         5      False
         6      False
         7      False
         8      False
         9      False
        10      False
        11      False
        12      False
        13      False
        14      False
        15      False
        16      False
        17      False
        18      False
        19      False
        20      False
        21      False
        22      False
        23      False
        24      False
        25      False
        26      False
        27      False
        28      False
        29      False
        ...
      2892      False
```

```
2893    False
2894    False
2895    False
2896    False
2897    False
2898    False
2899    False
2900    False
2901    False
2902    False
2903    False
2904    False
2905    False
2906    False
2907    False
2908    False
2909    False
2910    False
2911    False
2912    False
2913    False
2914    False
2915    False
2916    False
2917    False
2918    False
2919    False
2920    False
2921    False
dtype: bool
```

```
In [33]: # Try to pass location as the duplication column and see what happens
my_df.duplicated("location").value_counts()
```

```
# what is value_counts?
```

```
Out[33]: True      2920
False      2
dtype: int64
```

```
In [34]: # now try it with both location and weather columns
my_df.duplicated(["location", "weather"]).value_counts()
```

```
Out[34]: True      2912
False      10
dtype: int64
```

```
In [39]: # You can also remove duplicates using .drop_duplicates
len(my_df)
```

Out[39]: 2922

```
In [38]: # there are no duplicates of all the values
len(my_df.drop_duplicates())
```

Out[38]: 2922

```
In [40]: # You want to get the different weather that each location would get
# without any duplication
len(my_df.drop_duplicates(subset=["location", "weather"]))
```

Out[40]: 10

```
In [41]: # The first entry will be selected for each location and weather combination
my_df.drop_duplicates(subset=["location", "weather"])
```

```
Out[41]:
```

	location	date	precipitation	temp_max	temp_min	wind	weather
0	Seattle	2012-01-01	0.0	12.8	5.0	4.7	drizzle
1	Seattle	2012-01-02	10.9	10.6	2.8	4.5	rain
7	Seattle	2012-01-08	0.0	10.0	2.8	2.0	sun
13	Seattle	2012-01-14	4.1	4.4	0.6	5.3	snow
192	Seattle	2012-07-11	0.0	27.8	13.3	2.9	fog
1461	New York	2012-01-01	1.8	10.0	3.3	5.1	rain
1462	New York	2012-01-02	0.0	10.0	0.6	8.7	sun
1470	New York	2012-01-10	0.0	8.9	-1.1	5.5	drizzle
1473	New York	2012-01-13	0.0	10.0	-1.7	11.4	snow
1539	New York	2012-03-19	0.0	17.8	7.2	3.1	fog

```
In [43]: # you can even show the location and whather only using indexing
my_df.drop_duplicates(subset=["location", "weather"])[["location", "weather"]]
```

```
Out[43]:
```

	location	weather
0	Seattle	drizzle
1	Seattle	rain
7	Seattle	sun
13	Seattle	snow
192	Seattle	fog
1461	New York	rain
1462	New York	sun
1470	New York	drizzle
1473	New York	snow
1539	New York	fog